

```
In [ ]: ▶ import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import scale
import sklearn.linear_model as lm
import statsmodels.formula.api as smf
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
```

```
In [ ]: ▶ from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
In [ ]: ▶ adv = pd.read_csv('/content/gdrive/MyDrive/ISLR/Advertising.csv', usecols=[1, 2, 3, 4])
```

```
In [ ]: ▶ adv.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    radio       200 non-null    float64
2    newspaper   200 non-null    float64
3    sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [ ]: ▶ credit = pd.read_csv('/content/gdrive/MyDrive/ISLR/Credit.csv')
```

```
In [ ]: credit.head()
```

Out[17]:

	Unnamed: 0	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	Ethnicity	Balance
0	1	14.891	3606	283	2	34	11	Male	No	Yes	Caucasian	333
1	2	106.025	6645	483	3	82	15	Female	Yes	Yes	Asian	903
2	3	104.593	7075	514	4	71	11	Male	No	No	Asian	580
3	4	148.924	9504	681	3	36	11	Female	No	No	Asian	964
4	5	55.882	4897	357	2	68	16	Male	No	Yes	Caucasian	331

```
In [ ]: credit['Student2']=credit.Student.map({'No':0, 'Yes':1})
credit.head()
```

Out[18]:

	Unnamed: 0	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	Ethnicity	Balance	Student2
0	1	14.891	3606	283	2	34	11	Male	No	Yes	Caucasian	333	0
1	2	106.025	6645	483	3	82	15	Female	Yes	Yes	Asian	903	1
2	3	104.593	7075	514	4	71	11	Male	No	No	Asian	580	0
3	4	148.924	9504	681	3	36	11	Female	No	No	Asian	964	0
4	5	55.882	4897	357	2	68	16	Male	No	Yes	Caucasian	331	0






```
In [ ]: ▶ ##Least squares fit
sns.regplot(adv.TV,adv.sales,order=1)    ##check ploynomila or not
```

/usr/local/lib/python3.8/dist-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variables as keyword arguments: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x7f52ccf1f0a0>
```



```
In [ ]: ▶ lr = lm.LinearRegression().
X = scale(adv.TV,with_mean=True,with_std=False).reshape(-1,1)
Y = adv.sales
lr.fit(X,Y)
print(lr.intercept_)
print(lr.coef_)
```

```
14.0225
[0.04753664]
```

```
In [ ]: ▶ RSS = np.sum((lr.intercept_+lr.coef_*X-Y.values.reshape(-1,1))**2)
```

```
In [ ]: ▶ RSS
```

```
Out[25]: 2102.5305831313512
```

```
In [ ]: ▶ lm_fit1 = smf.ols('sales ~ TV', adv).fit()  
lm_fit1.summary().tables[1]
```

```
Out[26]:
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	7.0326	0.458	15.360	0.000	6.130	7.935
TV	0.0475	0.003	17.668	0.000	0.042	0.053

```
In [ ]: ▶ adv.head(3)
```

```
Out[27]:
```

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3

```
In [ ]: ▶ lm_fit2 = smf.ols('sales ~ newspaper', adv).fit()  
lm_fit2.summary().tables[1]
```

```
Out[28]:
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	12.3514	0.621	19.876	0.000	11.126	13.577
newspaper	0.0547	0.017	3.300	0.001	0.022	0.087

```
In [ ]: ▶ ##Multiple Linear Regression  
lm_fit3 = smf.ols('sales ~ TV+radio+newspaper', adv).fit()  
lm_fit3.summary().tables[1]
```

```
Out[29]:
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.9389	0.312	9.422	0.000	2.324	3.554
TV	0.0458	0.001	32.809	0.000	0.043	0.049
radio	0.1885	0.009	21.893	0.000	0.172	0.206
newspaper	-0.0010	0.006	-0.177	0.860	-0.013	0.011

```
In [ ]: ► ##Correlation Matrix
adv.corr()
```

```
Out[30]:
```

	TV	radio	newspaper	sales
TV	1.000000	0.054809	0.056648	0.782224
radio	0.054809	1.000000	0.354104	0.576223
newspaper	0.056648	0.354104	1.000000	0.228299
sales	0.782224	0.576223	0.228299	1.000000

```
In [ ]: ► ##RSS for Multiple regression
##https://github.com/JWarmenhoven/ISLR-python/blob/master/Notebooks/Chapter%203.ipynb
```

```
lr = lm.LinearRegression()
```

```
X = adv[['radio', 'TV']]
```

```
Y = adv.sales
```

```
lr.fit(X,Y)
```

```
print(lr.intercept_)
```

```
print(lr.coef_)
```

```
2.9210999124051398
```

```
[0.18799423 0.04575482]
```

```
In [ ]: ► ##Other Considerations in the Regression Model
credit.head(2)
```

```
Out[32]:
```

	Unnamed: 0	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	Ethnicity	Balance	Student2
0	1	14.891	3606	283	2	34	11	Male	No	Yes	Caucasian	333	0
1	2	106.025	6645	483	3	82	15	Female	Yes	Yes	Asian	903	1


```
In [ ]: sns.pairplot(credit[['Balance', 'Age', 'Cards', 'Education', 'Income', 'Limit', 'Rating']])
```

```
Out[33]: <seaborn.axisgrid.PairGrid at 0x7f52cd3e14f0>
```









In [ ]:

lm\_fit1 = smf.ols('Balance~Ethnicity',credit).fit()  
lm\_fit1.summary().tables[1]

Out[35]:

	coef	std err	t	P> t	[0.025	0.975]
Intercept	531.0000	46.319	11.464	0.000	439.939	622.061
Ethnicity[T.Asian]	-18.6863	65.021	-0.287	0.774	-146.515	109.142
Ethnicity[T.Caucasian]	-12.5025	56.681	-0.221	0.826	-123.935	98.930

In [ ]:

adv.head(2)

Out[36]:

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4

In [ ]:

##Interaction effect  
lm\_fit2 = smf.ols('sales~TV+radio+TV\*radio',adv).fit()  
lm\_fit2.summary().tables[1]


Out[37]:

	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.7502	0.248	27.233	0.000	6.261	7.239
TV	0.0191	0.002	12.699	0.000	0.016	0.022
radio	0.0289	0.009	3.241	0.001	0.011	0.046
TV:radio	0.0011	5.24e-05	20.727	0.000	0.001	0.001

```
In [ ]: credit.head(2)
```

Out[38]:

	Unnamed: 0	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	Ethnicity	Balance	Student2
0	1	14.891	3606	283	2	34	11	Male	No	Yes	Caucasian	333	0
1	2	106.025	6645	483	3	82	15	Female	Yes	Yes	Asian	903	1

```
In [ ]:  ##Interaction between qualitative and quantative variables  
lm_fit3 = smf.ols('Balance~Income',credit).fit()  
print(lm_fit3.summary())  
lm_fit4 = smf.ols('Balance~Income+Student2+Income*Student2',credit).fit()  
print(lm_fit4.summary())
```

# OLS Regression Results

```

=====
Dep. Variable:          Balance    R-squared:                0.215
Model:                  OLS        Adj. R-squared:            0.213
Method:                 Least Squares    F-statistic:              109.0
Date:                  Thu, 29 Dec 2022    Prob (F-statistic):       1.03e-22
Time:                  05:50:06    Log-Likelihood:           -2970.9
No. Observations:      400        AIC:                      5946.
Df Residuals:          398        BIC:                      5954.
Df Model:               1
Covariance Type:       nonrobust
=====

```

```

=====
               coef    std err          t      P>|t|      [0.025    0.975]
-----
Intercept    246.5148    33.199      7.425    0.000    181.247    311.783
Income        6.0484     0.579    10.440    0.000     4.909     7.187
=====

```

```

=====
Omnibus:                 42.505    Durbin-Watson:           1.951
Prob(Omnibus):           0.000    Jarque-Bera (JB):        20.975
Skew:                    0.384    Prob(JB):                 2.79e-05
Kurtosis:                 2.182    Cond. No.                  93.3
=====

```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

# OLS Regression Results

```

=====
Dep. Variable:          Balance    R-squared:                0.280
Model:                  OLS        Adj. R-squared:            0.274
Method:                 Least Squares    F-statistic:              51.30
Date:                  Thu, 29 Dec 2022    Prob (F-statistic):       4.94e-28
Time:                  05:50:06    Log-Likelihood:           -2953.7
No. Observations:      400        AIC:                      5915.
Df Residuals:          396        BIC:                      5931.
Df Model:               3
Covariance Type:       nonrobust
=====

```

```

=====
               coef    std err          t      P>|t|      [0.025    0.975]
-----
Intercept    200.6232    33.698      5.953    0.000    134.373    266.873
Income        6.2182     0.592    10.502    0.000     5.054     7.382
Student2     476.6758    104.351     4.568    0.000    271.524    681.827
Income:Student2  -1.9992     1.731     -1.155    0.249    -5.403     1.404
=====

```

```

=====
Omnibus:                 107.788    Durbin-Watson:           1.952
Prob(Omnibus):           0.000    Jarque-Bera (JB):        22.158
=====

```

Skew:	0.228	Prob(JB):	1.54e-05
Kurtosis:	1.941	Cond. No.	309.
=====			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [ ]: ##Non-Linear relationships
plt.scatter(Auto.horsepower, Auto.mpg)
sns.regplot(Auto.horsepower, Auto.mpg,order=1,label='linear',scatter=False,color='red')
sns.regplot(Auto.horsepower, Auto.mpg,order=2,label='order2',scatter=False,color='orange')
sns.regplot(Auto.horsepower, Auto.mpg,order=5,label='order5',scatter=False,color='green')
plt.show()
```

/usr/local/lib/python3.8/dist-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


warnings.warn(

/usr/local/lib/python3.8/dist-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

/usr/local/lib/python3.8/dist-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



```
In [ ]: ▶ Auto['horsepower2'] = Auto.horsepower**2
Auto.head(2)
```

```
Out[41]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name	horsepower2
0	18.0	8	307.0	130.0	3504	12.0	70	1	chevrolet chevelle malibu	16900.0
1	15.0	8	350.0	165.0	3693	11.5	70	1	buick skylark 320	27225.0

```
In [ ]: ▶ lm_fit4 = smf.ols('mpg~horsepower+horsepower2',Auto).fit()
lm_fit4.summary().tables[1]
```

```
Out[42]:
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	56.9001	1.800	31.604	0.000	53.360	60.440
horsepower	-0.4662	0.031	-14.978	0.000	-0.527	-0.405
horsepower2	0.0012	0.000	10.080	0.000	0.001	0.001

```
In [ ]: ▶ ##find out the oder
lr = lm.LinearRegression()
X = Auto.horsepower.values.reshape(-1,1)
Y = Auto.mpg
lr.fit(X,Y)
Auto['mpg_pred'] = lr.predict(X)
Auto['res'] = Auto.mpg-Auto.mpg_pred

lr1 = lm.LinearRegression()
X1 = Auto[['horsepower', 'horsepower2']]
Y = Auto.mpg
lr1.fit(X1,Y)
Auto['mpg_pred1'] = lr1.predict(X1)
Auto['res1'] = Auto.mpg-Auto.mpg_pred1
```



In [ ]: ▶

```
fig, (ax1,ax2) = plt.subplots(1,2, figsize=(12,5))

# Left plot
sns.regplot(Auto.mpg_pred, Auto.res,lowess=True,
            ax=ax1, line_kws={'color':'r', 'lw':1},
            scatter_kws={'facecolors':'None', 'edgecolors':'k', 'alpha':0.5})
ax1.hlines(0,xmin=ax1.xaxis.get_data_interval()[0],
          xmax=ax1.xaxis.get_data_interval()[1], linestyle='dotted')
ax1.set_title('Residual Plot for Linear Fit')

# Right plot
sns.regplot(Auto.mpg_pred1, Auto.res1, lowess=True,
            line_kws={'color':'r', 'lw':1}, ax=ax2,
            scatter_kws={'facecolors':'None', 'edgecolors':'k', 'alpha':0.5})
ax2.hlines(0,xmin=ax2.xaxis.get_data_interval()[0],
          xmax=ax2.xaxis.get_data_interval()[1], linestyle='dotted')
ax2.set_title('Residual Plot for Quadratic Fit')


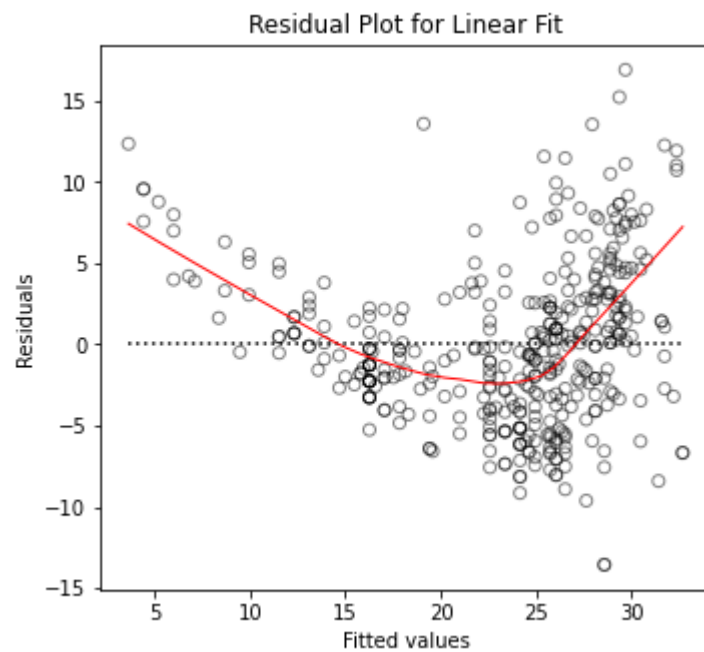
for ax in fig.axes:
    ax.set_xlabel('Fitted values')
    ax.set_ylabel('Residuals')
```

/usr/local/lib/python3.8/dist-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

/usr/local/lib/python3.8/dist-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



```
In [ ]: ##VIF-Variance Inflation Factor
est_Age = smf.ols('Age ~ Rating + Limit', credit).fit()
est_Rating = smf.ols('Rating ~ Age + Limit', credit).fit()
est_Limit = smf.ols('Limit ~ Age + Rating', credit).fit()

print(1/(1-est_Age.rsquared))
print(1/(1-est_Rating.rsquared))
print(1/(1-est_Limit.rsquared))
```

```
1.0113846860681328
160.66830095856935
160.59287978597942
```

```
In [ ]: ##LAB Assinments
df = pd.read_csv('/content/gdrive/MyDrive/ISLR/Boston.csv')
df.head(2)
```

```
Out[4]:
```

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.9	4.98	24.0
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.9	9.14	21.6

```
In [ ]: lm_fit = smf.ols('medv~lstat',df).fit()
print(lm_fit.params)
print()
print(lm_fit.conf_int())
print()
print(lm_fit.summary().tables[1])
```

```
Intercept    34.553841
lstat        -0.950049
dtype: float64
```

```
           0           1
Intercept  33.448457  35.659225
lstat      -1.026148  -0.873951
```

```
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    34.5538      0.563      61.415      0.000      33.448      35.659
lstat        -0.9500      0.039     -24.528      0.000      -1.026      -0.874
=====
```

```
In [ ]: lm_fit.predict(pd.DataFrame({'lstat':[5,10,15]}))
```

```
Out[15]: 0    29.803594
         1    25.053347
         2    20.303101
dtype: float64
```

```
In [ ]: ##OR
lr = lm.LinearRegression()
X = df.lstat.values.reshape(-1,1)
y = df.medv
lr.fit(X,y)
x_test = pd.DataFrame([5,10,15])
y_pred = lr.predict(x_test)
print(y_pred)
```


```
[29.80359411 25.05334734 20.30310057]
```

```
In [ ]: ▶ plt.scatter(df.lstat, df.medv)
sns.regplot(df.lstat, df.medv, order=1, label='linear', scatter=False, color='red')
```

/usr/local/lib/python3.8/dist-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variables as keyword arguments: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8f65c0deb0>
```



```
In [ ]: ▶ influence = lm_fit.get_influence()
leverage = influence.hat_matrix_diag
np.argmax(leverage)
```

```
Out[17]: 374
```


```
In [ ]: ▶ lm_fit = smf.ols('medv~lstat+age',df).fit()
lm_fit.summary().tables[1]
```

```
Out[7]:
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	33.2228	0.731	45.458	0.000	31.787	34.659
lstat	-1.0321	0.048	-21.416	0.000	-1.127	-0.937
age	0.0345	0.012	2.826	0.005	0.011	0.059

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   506 non-null   int64
1   crim         506 non-null   float64
2   zn           506 non-null   float64
3   indus        506 non-null   float64
4   chas         506 non-null   int64
5   nox          506 non-null   float64
6   rm           506 non-null   float64
7   age          506 non-null   float64
8   dis          506 non-null   float64
9   rad          506 non-null   int64
10  tax          506 non-null   int64
11  ptratio      506 non-null   float64
12  black        506 non-null   float64
13  lstat        506 non-null   float64
14  medv         506 non-null   float64
dtypes: float64(11), int64(4)
memory usage: 59.4 KB
```

```
In [ ]:  #all_columns = "+".join(Boston.columns.difference(["medv"]))
#my_formula = "medv~" + all_columns
#lm = smf.ols(my_formula, data=Boston).fit()
lm_fit = smf.ols('medv~crim+zn+indus+chas+nox+rm+age+dis+rad+tax+ptratio+lstat',df).fit()
lm_fit.summary().tables[1]
```

Out[10]:

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	41.6173	4.936	8.431	0.000	31.919	51.316
<b>crim</b>	-0.1214	0.033	-3.678	0.000	-0.186	-0.057
<b>zn</b>	0.0470	0.014	3.384	0.001	0.020	0.074
<b>indus</b>	0.0135	0.062	0.217	0.829	-0.109	0.136
<b>chas</b>	2.8400	0.870	3.264	0.001	1.131	4.549
<b>nox</b>	-18.7580	3.851	-4.870	0.000	-26.325	-11.191
<b>rm</b>	3.6581	0.420	8.705	0.000	2.832	4.484
<b>age</b>	0.0036	0.013	0.271	0.787	-0.023	0.030
<b>dis</b>	-1.4908	0.202	-7.394	0.000	-1.887	-1.095
<b>rad</b>	0.2894	0.067	4.325	0.000	0.158	0.421
<b>tax</b>	-0.0127	0.004	-3.337	0.001	-0.020	-0.005
<b>ptratio</b>	-0.9375	0.132	-7.091	0.000	-1.197	-0.678
<b>lstat</b>	-0.5520	0.051	-10.897	0.000	-0.652	-0.452

```
In [ ]: print(df.iloc[:,1:])
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax \
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222
..	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273

	ptratio	black	lstat	medv	intercept
0	15.3	396.90	4.98	24.0	1
1	17.8	396.90	9.14	21.6	1
2	17.8	392.83	4.03	34.7	1
3	18.7	394.63	2.94	33.4	1
4	18.7	396.90	5.33	36.2	1
..	...	...	...	...	...
501	21.0	391.99	9.67	22.4	1
502	21.0	396.90	9.08	20.6	1
503	21.0	396.90	5.64	23.9	1
504	21.0	393.45	6.48	22.0	1
505	21.0	396.90	7.88	11.9	1

[506 rows x 15 columns]

In [ ]: ▶

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant
X = add_constant(df)
#X = X.iloc[:,1:]

vif_df = pd.DataFrame()
vif_df["feature"] = X.columns

# calculating VIF for each feature
vif_df["VIF"] = [variance_inflation_factor(X.values, i)
                 for i in range(len(X.columns))]

#Print the dataframe with VIF values
print(vif_df)
```

	feature	VIF
0	crim	1.767486
1	zn	2.298459
2	indus	3.987181
3	chas	1.071168
4	nox	4.369093
5	rm	1.912532
6	age	3.088232
7	dis	3.954037
8	rad	7.445301
9	tax	9.002158
10	ptratio	1.797060
11	lstat	2.870777
12	intercept	535.526619

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

```
x = pd.concat(x[::order], 1)
```



```
In [ ]: lm_fit = smf.ols('medv~lstat+age+lstat*age',df).fit()
lm_fit.summary().tables[1]
```

```
Out[55]:
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	36.0885	1.470	24.553	0.000	33.201	38.976
lstat	-1.3921	0.167	-8.313	0.000	-1.721	-1.063
age	-0.0007	0.020	-0.036	0.971	-0.040	0.038
lstat:age	0.0042	0.002	2.244	0.025	0.001	0.008

```
In [ ]: lm_order1 = smf.ols('medv~ lstat', data=df).fit()
lm_order2 = smf.ols('medv~ lstat+ I(lstat ** 2.0)', data=df).fit()
print(lm_order2.summary().tables[1])
```

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	42.8620	0.872	49.149	0.000	41.149	44.575
lstat	-2.3328	0.124	-18.843	0.000	-2.576	-2.090
I(lstat ** 2.0)	0.0435	0.004	11.628	0.000	0.036	0.051

```
=====
```

```
In [ ]: import statsmodels.api as sm
table = sm.stats.anova_lm(lm_order1, lm_order2)
print(table)  ##order2 is a superior model
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	504.0	19472.381418	0.0	NaN	NaN	NaN
1	503.0	15347.243158	1.0	4125.13826	135.199822	7.630116e-28

```
In [ ]: ##OR
df['lstat2']=df.lstat**2
lm_fit = smf.ols('medv~lstat+lstat2',df).fit()
lm_fit.summary().tables[1]
```

```
Out[58]:
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	42.8620	0.872	49.149	0.000	41.149	44.575
lstat	-2.3328	0.124	-18.843	0.000	-2.576	-2.090
lstat2	0.0435	0.004	11.628	0.000	0.036	0.051

```
In [ ]: df = pd.read_csv('/content/gdrive/MyDrive/ISLR/carseats.csv')
df.head(2)
```

```
Out[23]:
```

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban	US
0	9.50	138	73	11	276	120	Bad	42	17	Yes	Yes
1	11.22	111	48	16	260	83	Good	65	10	Yes	Yes

```
In [ ]: lm_fit = smf.ols('Sales ~ Income + Advertising + Price + Age',df).fit()
print(lm_fit.summary().tables[1])
```

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	15.1829	0.777	19.542	0.000	13.656	16.710
Income	0.0108	0.004	2.664	0.008	0.003	0.019
Advertising	0.1203	0.017	7.078	0.000	0.087	0.154
Price	-0.0573	0.005	-11.932	0.000	-0.067	-0.048
Age	-0.0486	0.007	-6.956	0.000	-0.062	-0.035

```
=====
```

```
In [ ]: ▶ ShelfLoc_dummies = pd.get_dummies(df.ShelveLoc,prefix='ShelveLoc').iloc[:,1:]
ShelveLoc_dummies
df_dummy = pd.concat([df, ShelfLoc_dummies], axis=1)
df_dummy.head(2)
```

```
Out[31]:
```

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban	US	ShelveLoc_Good	ShelveLoc_Medium
0	9.50	138	73	11	276	120	Bad	42	17	Yes	Yes	0	0
1	11.22	111	48	16	260	83	Good	65	10	Yes	Yes	1	0

```
In [ ]: ▶ #https://github.com/qx0731/Sharing_ISL_python/blob/master/Chapter_3_sec_6.1_6.7.ipynb
```

```
In [ ]: ▶ lm_df_dummy = smf.ols('Sales ~ Income + Advertising + Price + Age + ShelveLoc_Good + ShelveLoc_Medium', data = df_dummy).fi
print(lm_df_dummy.summary().tables[1])
```

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	13.4006	0.545	24.575	0.000	12.329	14.473
Income	0.0136	0.003	4.891	0.000	0.008	0.019
Advertising	0.1057	0.012	9.076	0.000	0.083	0.129
Price	-0.0606	0.003	-18.436	0.000	-0.067	-0.054
Age	-0.0498	0.005	-10.401	0.000	-0.059	-0.040
ShelveLoc_Good	4.8756	0.230	21.175	0.000	4.423	5.328
ShelveLoc_Medium	2.0046	0.189	10.590	0.000	1.632	2.377

```
=====
```

```
In [ ]: ##OR
lm_df_w_dummy = smf.ols('Sales ~ Income + Advertising + Price + Age + C(ShelveLoc)', data = df).fit()
print(lm_df_w_dummy.summary().tables[1])
```

```
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept          13.4006         0.545      24.575      0.000      12.329      14.473
C(ShelveLoc)[T.Good]    4.8756         0.230      21.175      0.000         4.423         5.328
C(ShelveLoc)[T.Medium]  2.0046         0.189      10.590      0.000         1.632         2.377
Income              0.0136         0.003         4.891      0.000         0.008         0.019
Advertising          0.1057         0.012         9.076      0.000         0.083         0.129
Price              -0.0606         0.003     -18.436      0.000        -0.067        -0.054
Age                -0.0498         0.005     -10.401      0.000        -0.059        -0.040
=====
```

```
In [ ]: ## one_hot_encoded_data = pd.get_dummies(credit, columns = ['Married', 'Student', 'Gender'])
print(one_hot_encoded_data)
```

```
In [ ]: ## ##1. NA value, replace with mean, median or mode, quantitative then mean, qualitative mode or median
##2. covert X and Y into scale
###3. df.info( ), factor variable and continuous variable
##4. VIF and check which variables i should drop
##5. check linear or non linear
##6. do simple LR model then findout the factor variable whcih is significant
##7. those factor which is significant do onehotencode
###8. TRY to findout best model where all variables p values less than 0.05, R sq
##and adj R sq high and AIC and BIC low.
## this is only for traning model
```